CLINICAL QUANTUM ROADMAP • NHS LONG TERM CONDITIONS • SAMSKRIFT RESEARCH MONOGRAPH

Quantum-LTC: A Four-Era Clinical Quantum Optimisation and Simulation Roadmap Towards Fault-Tolerant Quantum Advantage for NHS Long Term Conditions

Arun Nadarasa[1,2,3]

[1]Innovation Fellowship Programme, Mid and South Essex NHS Foundation Trust, Basildon, Southend & Broomfield, UK
[2]International Social Prescribing Pharmacy Association (ISPPA), London, UK
[3]Primary Care & Integrated Neighbourhood Health Systems Research Group, London, UK

Portal & Live Space: arunnadarasa.github.io/quantum-ltc • hf.co/spaces/socialprescribing/quantum-ltc | **HF Model:** quantum-ltc-qltcmlm

Abstract. Managing multimorbid Long Term Conditions (LTCs) across primary care networks requires reconciling constrained clinic capacity, pre-winter high-risk stratification, polypharmacy safety, health-equity weighting, and decarbonisation targets [1, 2, 3]. Grounded in the **NHS North Central London Integrated Care Board (NCL ICB) Locally Commissioned Service (LCS) for Proactive, Personalised Care — Long Term Conditions (Version 6, 2026–2027)** [1], which unifies nine cardiometabolic, renal, hepatic, and respiratory conditions within a four-stage House-of-Care pathway [4, 5], we present **Quantum-LTC**: an end-to-end, cryptographically receipted clinical quantum optimisation and biophysical simulation roadmap spanning $N_{\text{logical}} \in [0, 1{,}000+]$ logical qubits. Crucially, we decouple *sparse discrete combinatorial scheduling*—where exact classical McCormick/Rosenberg mixed-integer linear programming (MILP) via HiGHS [6, 7, 8] dominates NISQ heuristics—from *strongly coupled open-system Lindbladian dynamics* and *metalloenzyme active-space quantum chemistry*, where full-density-matrix $\mathcal{O}(4^{n_{\text{spin}}})$ Krylov RAM walls, volume-law Matrix Product Density Operator (MPDO) bond growth [9], and $\mathcal{O}(1/\epsilon^2)$ Monte Carlo Wave Function (MCWF) trajectory variance [10] establish the earliest fault-tolerant crossover. Across three live **Guppy 1.0 → HUGR** jobs on **Quantinuum Nexus** `Helios-1E-lite` (6,144 emulator shots; 0 physical QPU shots), coherent Stage 1–4 multidisciplinary team (MDT) synthesis achieves $P(00) = 0.9961 \pm 0.0014$ versus 0.5020 ± 0.0110 under equal-unitary-depth siloed single-disease collapse, while an 8-rung mid-circuit measurement and reset (MCMR) recall ladder achieves $R^2 = 0.9979$ state preparation and measurement (SPAM) calibration linearity. On an 8-pathway 3-body cardiometabolic–renal–respiratory and Net-Zero dry-powder MART inhaler HUBO [11, 12], **Haiqu SDK** compresses $p = 10$ Linear-Ramp QAOA [13] depth from 223 → 26 (−88.34%), yielding 41.21% ± 1.54% raw ground-state probability (105.5× uniform random) and exact −23.500 ground energy. On a 128-variable NCL Call-and-Recall QUBO (**JijModeling/OpenJij**), raw simulated annealing (160/160 hard-constraint feasible, −113.85) plus 4-bit Kempe polish lifts pre-January High/Complex MDT recall from 14/18 (77.8%) to 18/18 (100.0%) within 0.80% (−134.57) of the 4.33 ms HiGHS MILP floor (−135.66). Biophysical, cloud QML, and governance modules integrate decoupled **HQS** `struqture-py/qoqo` 4-spin Lindbladians, **QProtSim** on `PDB 1N9V` [14], ancilla-complete `PDB 1O86` (`ACE`–lisinopril Zn^{2+}, 2.141 Å) FTQC estimates [15], **AlphaGenome Atlas**, 5-seed **Modal Cloud** `QLTCMLM` reservoir/polynomial controls, **Rhyme** (`rhyme-beta`) [16], **Text2WetLab** OT-2 checks, **NIST FIPS 204** `ML-DSA-65` FHIR R4 provenance [17], **Kashefi Primitives Atlas v0.4** Lie-algebra `g`-sim audits [18, 19], and live **TypeSafe AI System One (`jev-latest` / Jev 1.14)** safety auditing under statutory **NHS DCB0129** and **DCB0160** [20, 21, 22].

Epistemic Demarcation & Statutory NHS DCB0129 / DCB0160 Clinical Safety Disclosure. All experimental results and ablations are deterministically reproducible from **six SHA-256-sealed JSON receipts** in `results/` (`0e3216dd...`, `7212c3d9...`, `6ba83347...`, `cb1fec7c...`, `bd60742f...`, `a6b3854e...`) plus the 57-node knowledge graph (`9b4e62e0...`). Every cohort ($N = 2{,}400$), PCN schedule ($P = 32$), HUBO tensor, and phenomenological Lindbladian rate is synthetic (`phi_present = False`; **0 physical QPU hardware shots**; **0 wet-lab human specimens**); in particular, the 75% cubic ($k = 3$) cohort interaction weighting is a stylized non-linear stress-test rather than an empirical fit to UK primary-care EHR distributions (e.g., CPRD Aurum) or validated NHS scores such as the Electronic Frailty Index (**eFI/eFI2**) [23] and **Cambridge Multimorbidity Score (CMS)** [24]. Classical HiGHS MILP and polynomial/kernel baselines dominate all Era I–II tasks; fault-tolerant crossover is projected solely in Era III ($N_{\text{logical}} \in [38, 68]$) against exact $\mathcal{O}(4^{n_{\text{spin}}})$ density-matrix integration, volume-law MPDOs, and $\mathcal{O}(1/\epsilon^2)$ MCWF sampling, and Era IV ($N_{\text{logical}} \geq 100$) on dense frustrated $k \geq 4$ hypergraphs and QSVT/QAE tail risk. All six clinical hazards (`H-LTC-01`–`H-LTC-06`) remain strictly **OPEN** across both manufacturer-side **NHS DCB0129** [21] and deploying-organisation **NHS DCB0160** [22] governance under mandatory GP, Clinical Pharmacist, and Tier 2/3 Specialist Nurse sign-off.

1. Clinical Specification: The NHS NCL ICB Multimorbid LTC Service Model

Multimorbidity—the co-occurrence of two or more chronic conditions within an individual—is projected by the Health Foundation REAL Centre to expand substantially across England through 2040, with severe illness concentrated disproportionately in the most socioeconomically deprived quintile [2]. Traditional Quality and Outcomes Framework (QOF) single-disease templates frequently subject multimorbid patients to fragmented, duplicative clinic appointments, conflicting polypharmacy adjustments, and uncoordinated lifestyle advice [3, 4]. To overcome single-condition fragmentation across Barnet, Camden, Enfield, Haringey, and Islington, the **NHS North Central London Integrated Care Board (NCL ICB)** commissioned Version 6 of the *Locally Commissioned Service (LCS) for Proactive, Personalised Care — Long Term Conditions* (1 April 2026 – 31 March 2027) [1].

The NCL LTC LCS unifies **nine conditions** accounting for approximately 75% of modifiable cardiometabolic, renal, hepatic, and respiratory morbidity: **(1)** Cardiovascular Disease (coronary heart disease, stroke/TIA, peripheral arterial disease, heart failure), **(2)** Atrial Fibrillation, **(3)** Hypertension (excluding low-risk well-controlled cases, anchored by **NCL PCN Outcome #23**: achieving blood pressure $< 140/90$ mmHg [home $< 135/85$] for age < 80 or $< 150/90$ mmHg [home $< 145/85$] for age ≥ 80 across all five boroughs), **(4)** Type 2 Diabetes, **(5)** Type 1 Diabetes (mandated `High Risk + Complex`), **(6)** Chronic Kidney Disease (CKD), **(7)** Non-Alcoholic Fatty Liver Disease / Metabolic Dysfunction-Associated Steatotic Liver Disease (NAFLD/MASLD), **(8)** COPD, and **(9)** Adult and Children/Young People (CYP, aged 0–17) Asthma, which mandates a **48-hour post-attack review** (within two working days of an emergency or out-of-hours presentation) and a **6-week full review** by Tier 2/3 trained clinicians alongside short-acting β_2-agonist (SABA)-free pathways [1, 12].

Operationally, the specification embeds the four-stage **Year of Care / House of Care** cycle [4, 5]:

1. **Proactive Case Finding & Call-and-Recall Tension (Page 5 of LCS v6):** Practices run 3-yearly UCLPartners (UCLP) searches to close undiagnosed prevalence gaps and face an explicit scheduling dilemma between *Month-of-Birth recall* (which smooths staffing across 12 months but leaves high-risk patients exposed to winter admissions) and *Risk-Stratified recall* (`Low`, `Moderate`, `High`, `High Risk + Complex`), which prioritizes pre-winter proactive care (capable of reducing emergency admissions by up to 38% and GP contacts by 18% [5]) at the risk of seasonal MDT bottlenecks. In NHS GP EHRs (EMIS Web / SystmOne / Medicus), this 4-tier operational stratification complements validated primary-care frailty and multimorbidity instruments—specifically the cumulative-deficit **Electronic Frailty Index (eFI/eFI2)** [23] and the consultation/admission/mortality-weighted **Cambridge Multimorbidity Score (CMS)** [24]—by coupling longitudinal diagnostic deficits with acute physiological biomarkers (Outcome #23 BP, uACR/eGFR, HbA1c, post-attack CYP asthma) and linguistic/deprivation equity weights.

2. **Stage 1 (Check and Test) & Stage 2 (Re-Stratification & MDT Desktop Review):** A single 30-minute baseline appointment gathers phlebotomy (HbA1c, eGFR, lipids, LFTs/FIB-4, NT-proBNP), urine albumin-creatinine ratio (uACR), blood pressure, foot checks, and SNOMED CT `315594003` (`Interpreter needed`) access requirements, followed by Clinical Pharmacist / MDT Desktop Review for `High` and `High + Complex` patients.

3. **Stage 3 (PCSP Discussion) & Stage 4 (Continuity Follow-Up):** Two to four weeks after test results are shared, a 30-minute Coaching-for-Health appointment co-produces a Personalised Care and Support Plan (PCSP) integrating pharmacotherapy, Voluntary, Community and Social Enterprise (VCSE) social prescribing [25], and **NCL Green Plan Net Zero** switches (preventing carbon-intensive haemodialysis via early SGLT2i/ACEi protection [11] and replacing high-global-warming-potential propellant SABA inhalers with dry-powder Maintenance and Reliever Therapy [MART] [12]), followed by a clinician-continuity Stage 4 review evaluated via `collaboRATE` (≥ 25 Top-Score responses) and monthly prospective **PCN Weighted Payments** targeting the 20% most deprived areas (IMD Quintile 1) and minority ethnic communities [1].

2. The Four-Era Logical Qubit Roadmap: Decoupling Optimization from Simulation

A pervasive pitfall in translational quantum healthcare proposals is conflating the exponential operator-space growth of strongly entangled open quantum systems with sparse discrete combinatorial scheduling, which modern classical branch-and-cut solvers handle effortlessly via polynomial linearization [6, 7, 8]. To establish an intellectually rigorous trajectory as hardware progresses from noisy intermediate-scale quantum (NISQ) processors to modular trapped-ion fault tolerance [26, 27, 28], we partition clinical workloads across **four logical-qubit eras** (Table 1 and Figure 1).

2.1 Mathematical Scaling Separation: McCormick MILP, Trajectory/MPDO Baselines & FTQC

Consider first a general k-local clinical decision hypergraph over N binary interventions $x \in \{0,1\}^N$:

$$E_{\mathrm{HUBO}}(x) = \sum_{i=1}^{N} h_i x_i + \sum_{(i,j)\in E_2} J_{ij} x_i x_j + \sum_{(i,j,\ell)\in E_3} K_{ij\ell} x_i x_j x_\ell + \sum_{(i,j,\ell,m)\in E_4} M_{ij\ell m} x_i x_j x_\ell x_m. \tag{1}$$

Table 1: Four-Era Clinical Quantum Optimisation and Simulation Roadmap for NHS Long Term Conditions (N_{logical} Scaling). Physical qubit footprints follow $N_{\text{phys}} = (\lceil 1.5 N_{\text{logical}} \rceil + 12) \times 2d^2$ incorporating rotated surface-code tiles, lattice-surgery routing (1.5×), and a 12-tile magic-state factory ($p_{\text{phys}} = 10^{-3}$) [29, 30, 27]. Memory sizes use strict binary units (1 GiB = 2^{30} B, 1 TiB = 2^{40} B).

Roadmap Era	Logical Qubits	Physical Qubits	QEC & Hardware Regime	Target NHS NCL LTC Clinical & Biophysical Workloads	Classical vs. Quantum Verdict
Era I *Classical-Exact & NISQ Readiness*	0 logical (1–20 phys / cloud emul.)	1–98 (Helios / NISQ)	Uncorrected physical gates + 8-rung affine SPAM calibration	128-var PCN Call/Recall scheduling QUBO; 8-var 3-body multimorbid + Net-Zero MART HUBO; 4q Frobenius `QLTCMLM`	**Classical Dominates** (HiGHS MILP < 25 ms; 0 runtime speedup claimed)
Era II *Error-Detected Logical*	3–24 error-detected	10–192 (Helios / Sol QCCD)	Distance-2 $[\![8,3,2]\!]$ & $[\![4,2,2]\!]$ color codes + post-selection [27]	Shallow MCMR longitudinal recall combs (CYP Asthma 48h/6wk & BP #23); compressed $p \leq 3$–10 LR-QAOA circuits	**Classical MPS/MILP Faster** (Post-selection yield $p_{\text{acc}} \to 15.8\%$ at 24L, $d_{\text{circ}} = 24$)
Era III *Early Active FTQC Crossover*	24–72 active QEC	5,300–65,900 ($d = 9$–17 + T-factory)	Active multi-round decoding + magic-state cultivation / distillation	Bath-embedded 16–24 spin organ Lindbladians (38–54L); ancilla-complete `CAS(36e,26o)` (64–68L) `ACE` Zn^{2+}–`LPR` (`1086`)	**First Simulation Crossover** (Krylov > 1.875 TiB; volume-law MPDO & $\mathcal{O}(\epsilon^{-2})$ MCWF)
Era IV *Asymptotic ICS-Scale FTQC*	100–1,000+ active QEC	90,000–250,000+ ($d = 15$–17 modular)	Photonic-linked QCCD + parallel magic-state factories	Block-encoded QSVT organ twins (128L) [31]; $\mathcal{O}(1/\epsilon)$ QAE polypharmacy CVaR [32]; Quantum B&B on frustrated $k \geq 4$ HUBOs [33]	**Projected FTQC Advantage** (Advisory only under NHS DCB0129 & DCB0160 sign-off)

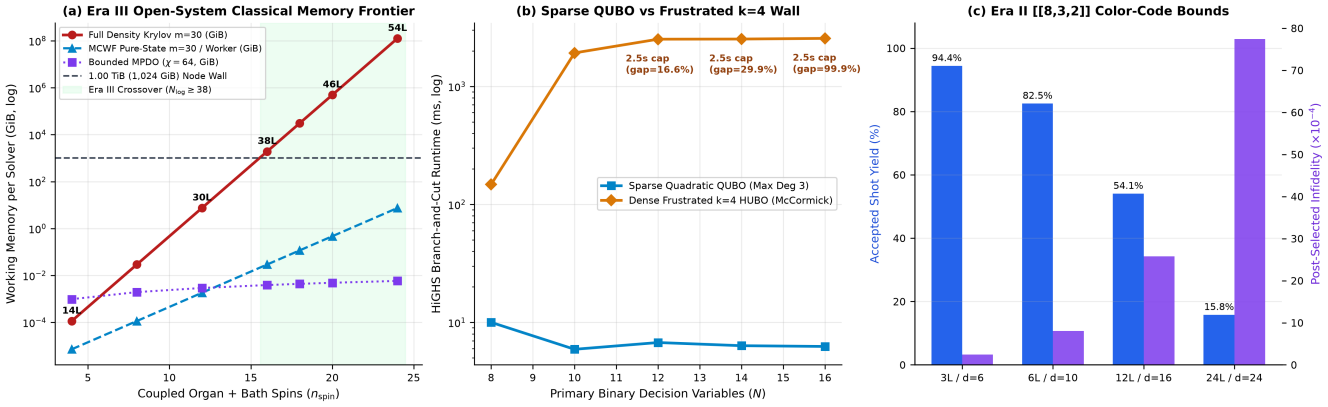


Figure 1: Quantitative Scaling Laws Across the Four-Era Logical Qubit Roadmap (`modal_roadmap_rhyme_qml_receipt.json` & `deep_review_remediation_receipt.json`). **(a)** Open-system n_{spin} Lindbladian memory in strict binary gibibytes (GiB = 2^{30} B): exact $m = 30$ full-density-matrix Krylov integration ($\dim \mathcal{L} = 4^{n_{\text{spin}}}$) crosses 1 TiB at $n_{\text{spin}} = 16$ (1,920 GiB = 1.875 TiB, 38 logical qubits) and reaches 480.0 TiB at $n_{\text{spin}} = 20$ (46 logical / 19,602 physical qubits at $d = 11$), contrasted against single-trajectory MCWF pure-state Krylov memory (30 MiB at $n_{\text{spin}} = 16$, 480 MiB at $n_{\text{spin}} = 20$, trading RAM for $\mathcal{O}(1/\epsilon^2)$ trajectory sampling [10]) and bounded ($\chi = 64$) vs. volume-law MPDOs [9]. **(b)** Empirical HiGHS branch-and-cut runtime comparing sparse quadratic QUBOs (max vertex degree 3; 6.26 ms at $N = 16$) against dense frustrated quartic ($k = 4$) HUBOs, which hit the 2.5 s cap from $N = 12$ upward with widening percentage MIP gaps (16.6% → 29.9% → 99.9%). **(c)** Analytical Era II $[\![8,3,2]\!]$ color-code post-selection shot yield $p_{\text{acc}} = (1 - p_{2q})^{N_{\text{loc}}}$ and post-selected logical infidelity $p_{\log,\text{eff}} = \min(1, 3.5 N_{\text{loc}} p_{2q}^2)$ at $p_{2q} = 1.2 \times 10^{-3}$.

Introducing continuous auxiliary variables $z_S = \prod_{i \in S} x_i \in [0,1]$ for each hyperedge S of order $|S| = k$ with exact McCormick linear inequalities [7]:

$$z_S \leq x_i \quad (\forall i \in S), \qquad z_S \geq \sum_{i \in S} x_i - (k-1), \qquad z_S \geq 0, \tag{2}$$

reduces Eq. (1) to an exact MILP with N binary variables, $|E_{\geq 2}|$ continuous variables, and $\sum_S (|S| + 1)$ structural linear constraint rows (since $z_S \geq 0$ is enforced directly via the solver box bound $z_S \in [0,1]$, yielding $5 \times 1{,}820 = 9{,}100$ rows at $N = 16, k = 4$). When the interaction hypergraph is sparse, dual-simplex + Gomory/MIR cutting planes in HiGHS [8] close the root LP gap in milliseconds. Under minimization ($\min_x E_{\text{HUBO}}(x)$) on dense frustrated instances containing competing positive and negative couplings, however, the fractional midpoint $x_i = 1/2$ satisfies $\sum_{i \in S} x_i - (k-1) = 1 - k/2 \leq 0$ for all $k \geq 2$ (strictly negative for $k \geq 3$): the LP relaxation simultaneously sets $z_S = 0$ freely on *positive-cost* (repulsive/penalty) hyperedges ($M_S > 0$, zeroing out frustration penalties even

though $\prod_{i\in S} x_i = 2^{-k} > 0$) while setting $z_S = 1/2 \gg 2^{-k}$ on *negative-cost* (attractive/reward) hyperedges ($M_S < 0$, over-collecting rewards by 8× at $k = 4$). Empirically (Figure 1b), while sparse quadratic QUBOs (maximum graph vertex degree 3) solve in 6.26 ms at $N = 16$, dense frustrated $\binom{N}{4}$ quartic HUBOs jump from 147.3 ms ($N = 8$, 70 terms) and 1,920.3 ms ($N = 10$, 210 terms) to hit the 2.5 s branch-and-cut cap at $N = 12, 14, 16$ (1,820 terms; residual MIP integrality gap 16.6% → 29.9% → 99.9%), motivating Era IV quantum branch-and-bound ($\mathcal{O}(\sqrt{T_{\text{tree}}})$) [33].

Conversely, consider open-system dynamics generated by a Gorini–Kossakowski–Sudarshan–Lindblad (GKSL) master equation over n_{spin} two-level degrees of freedom,

$$\frac{d\rho}{dt} = \mathcal{L}(\rho) = -i[H_{\text{organ}}, \rho] + \sum_{\mu} \gamma_\mu \left(L_\mu \rho L_\mu^\dagger - \frac{1}{2}\{L_\mu^\dagger L_\mu, \rho\} \right). \tag{3}$$

Mathematically, Eq. (3) with constant rates $\gamma_\mu \geq 0$ generates a strictly Markovian (CP-divisible) semigroup on the full n_{spin} register; non-Markovian memory-kernel dynamics on n_{organ} physiological compartments arise when $n_{\text{spin}} = n_{\text{organ}} + n_{\text{bath}}$ embeds explicit environmental/neurohumoral bath pseudomodes that are subsequently traced out ($\rho_{\text{organ}}(t) = \text{Tr}_{\text{bath}}\rho(t)$). Benchmarking quantum simulation fairly requires comparing against **three distinct classical open-system paradigms** (Figure 1a):

1. **Exact Full-Density-Matrix Krylov Integration ($\mathcal{O}(4^{n_{\text{spin}}})$):** Storing $m = 30$ double-precision complex vectors (16 B/entry) in $\mathbb{C}^{4^{n_{\text{spin}}}}$ requires $30 \times 16 \times 4^{n_{\text{spin}}}$ bytes: in strict binary units, 7.50 GiB at $n_{\text{spin}} = 12$, crossing 1 TiB at $n_{\text{spin}} = 16$ (**1,920** GiB = **1.875** TiB), **30.00** TiB at $n_{\text{spin}} = 18$, and **480.00** TiB at $n_{\text{spin}} = 20$ (527.77 TB decimal SI).

2. **Monte Carlo Wave Function (MCWF) Trajectory Unraveling [10]:** Propagating pure statevectors $|\psi(t)\rangle \in \mathbb{C}^{2^{n_{\text{spin}}}}$ interleaved with stochastic jumps reduces working memory to $\mathcal{O}(2^{n_{\text{spin}}})$—only 1.00 MiB (30.0 MiB for $m = 30$ Krylov) at $n_{\text{spin}} = 16$ and 16.00 MiB (480.0 MiB Krylov) at $n_{\text{spin}} = 20$—eliminating the single-node RAM wall at the expense of $\mathcal{O}(1/\epsilon^2)$ Monte Carlo trajectory variance ($\approx 10^6$ independent trajectories per parameter point to resolve $\epsilon = 10^{-3}$ steady-state shifts or rare adverse-event tails, compared with coherent $\mathcal{O}(1/\epsilon)$ QSVT/QAE scaling).

3. **Matrix Product Density Operators (MPDOs) [9]:** Locally purified tensor networks compress 1D/area-law dissipative chains into $\mathcal{O}(n_{\text{spin}}\chi^2)$ memory (≤ 5.0 MiB at $\chi = 64, n_{\text{spin}} = 20$), remaining faster than Era II–III hardware whenever operator-space entanglement stays bounded, but encountering exponential bond growth ($\chi \sim 2^{n_{\text{spin}}}$) when non-planar multi-organ crosstalk and strongly coupled bath pseudomodes generate volume-law operator entanglement.

Under Quantum Singular Value Transformation (QSVT) block-encoding [31], an $n_{\text{spin}} = 20$ bath-embedded Lindbladian requires $N_{\text{logical}} = 2n_{\text{spin}} + 6 = 46$ active logical qubits (40 Stinespring-dilated system/purification qubits, 4 linear-combination-of-unitaries [LCU] prepare/select ancillas indexing the jump operators L_μ, and 2 QSVT signal-processing ancillas). At an assumed T-gate budget of 7.2×10^4 ($p_{\text{phys}} = 10^{-3}$), rotated surface-code distance $d = 11$ ($2d^2 = 242$ physical qubits/tile) under

$$N_{\text{phys}}(N_{\text{logical}}, d) = (\lceil 1.5\, N_{\text{logical}} \rceil + N_{\text{factory}}) \times 2d^2, \qquad N_{\text{factory}} = 12, \tag{4}$$

yields $(69 + 12) \times 242 =$ **19,602** physical qubits, pinpointing **Era III (**38–68 **active logical qubits)** as the earliest crossover against exact Liouvillian and volume-law MPDO/chemistry regimes.

In **Era II (**3–24 **error-detected logical qubits)**, $B = N_{\text{logical}}/3$ blocks of the $[\![8, 3, 2]\!]$ color code ($n_{\text{phys}} = 8B$) [27] accrue $N_{\text{loc}} = 8B\, d_{\text{circ}} = n_{\text{phys}} d_{\text{circ}}$ effective qubit-layer fault locations across circuit depth d_{circ} (treating $p_{2q} = 1.2 \times 10^{-3}$ as an effective per-qubit-per-layer fault rate folding two-qubit gate, syndrome-extraction, and idling errors). Trivial-syndrome post-selection yields survival $p_{\text{acc}} = (1 - p_{2q})^{N_{\text{loc}}}$ and undetected weight-≥ 2 logical infidelity $p_{\text{log,eff}} = \min(1, 3.5 N_{\text{loc}} p_{2q}^2)$: at 3L ($B = 1, d_{\text{circ}} = 6, N_{\text{loc}} = 48$), $p_{\text{acc}} = 94.40\%$ and $p_{\text{log,eff}} = 2.42 \times 10^{-4}$, whereas at 24L ($B = 8, d_{\text{circ}} = 24, N_{\text{loc}} = 1{,}536$), survival drops to 15.81% (6.32× shot overhead; Figure 1c), necessitating active Era III syndrome decoding beyond $N_{\text{logical}} \approx 24$.

3. Live Trapped-Ion Emulation & Combinatorial Optimisation (Era I–II)

3.1 Live Quantinuum Nexus `Helios-1E-lite` Guppy 1.0 → HUGR Execution

We compiled three clinical verification circuits in **Guppy 1.0** directly to Quantinuum's Hierarchical Unified Graph Representation (**HUGR**) and executed them live on **Quantinuum Nexus `Helios-1E-lite`** (`QSystemErrorModel("alpha")`, project `Quantum-LTC-NCL`, ID `69eb4683-8de0-4dc9-bd22-ce6d98cafc7f`; $2{,}048 \times 3 = 6{,}144$ live cloud emulator shots; 0 physical QPU shots; Figure 2a–b):

1. **Equal-Unitary-Depth House-of-Care MDT Synthesis vs. Siloed Collapse** (Job UUID `5da15d5a-c48e-4a11-9bb3-6c0d8854131c`, HUGR `7,456 B`): Evaluates two 2-qubit arms simultaneously within a 4-qubit register. Acting *right-to-left* on $|00\rangle$ (so $(H \otimes I)$ and CX_{01} prepare Stage 1 Bell entanglement $|\Phi^+\rangle$ before local clinical-axis rotations and Stage 2 SWAP handoff), Arm A (q_0, q_1) applies $U_{\text{MDT}} = (H \otimes I)\, \text{CX}_{01}\, (R_z(-\theta_z) \otimes$

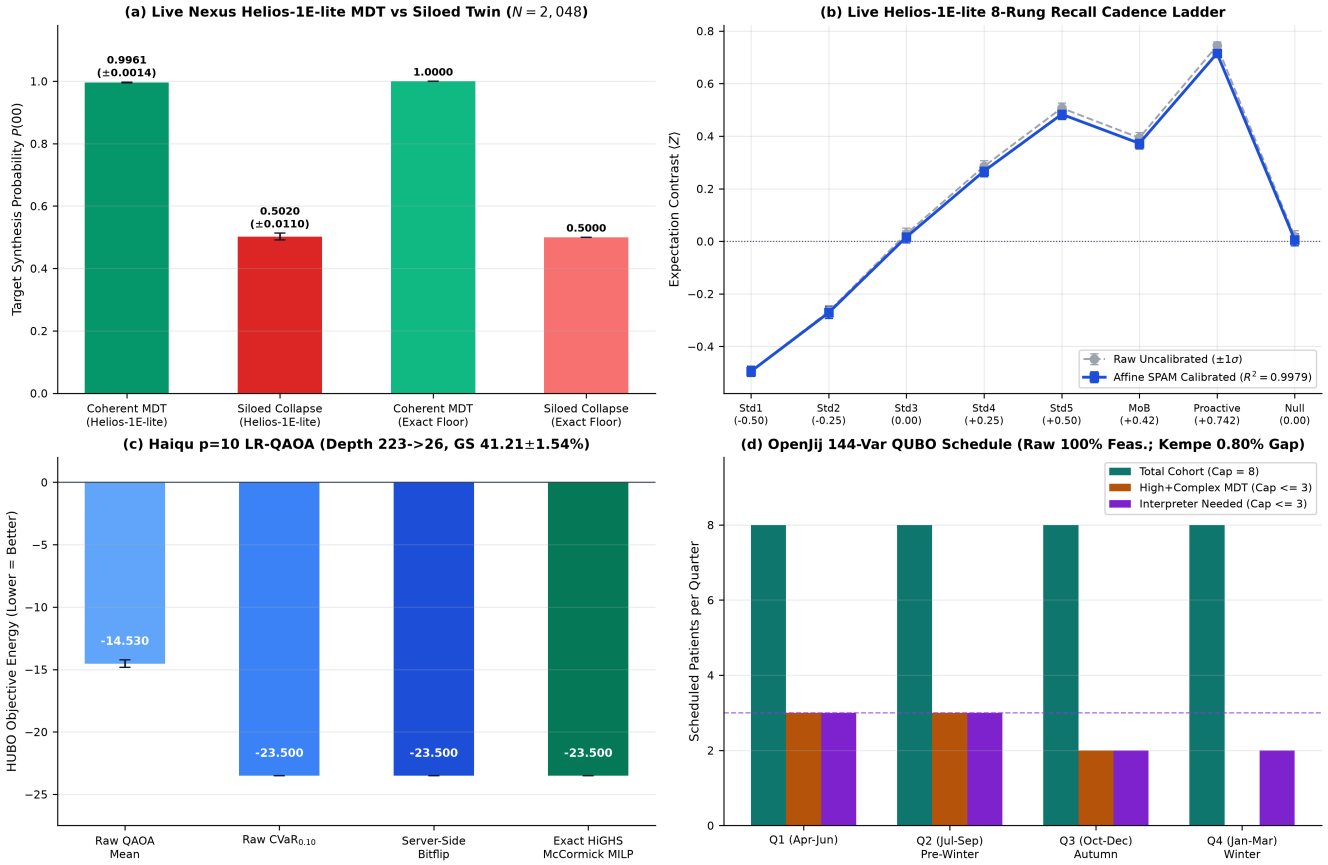


Figure 2: Live Quantinuum Nexus `Helios-1E-lite`, Haiqu 3-Body HUBO, and OpenJij NCL Call-and-Recall Benchmarks (`nexus_helios_ltc_receipt.json`, `haiqu_jij_ltc_receipt.json`, & `deep_review_remediation_receipt.json`; error bars denote $\pm1\sigma$ binomial/shot SE). **(a)** Equal-unitary-depth Stage 1–4 MDT synthesis ($P(00) = 0.9961 \pm 0.0014$) versus post-SWAP siloed collapse (0.5020 ± 0.0110). **(b)** 8-rung mid-circuit recall ladder ($R^2 = 0.9979$) across five SPAM standards $\{-0.50, -0.25, 0.00, +0.25, +0.50\}$ and three clinical recall profiles. **(c)** Haiqu SDK 8-qubit $p = 10$ LR-QAOA 3-body HUBO showing raw QAOA mean (-14.530 ± 0.292, raw GS hit rate $41.21\% \pm 1.54\%$), $\mathrm{CVaR}_{0.10}$ (-23.500), server-side bitflip postprocessing (-23.500), and exact HiGHS MILP floor (-23.500). **(d)** OpenJij + 4-bit Kempe-swap 128-variable NCL PCN Call-and-Recall schedule satisfying exact 8/8/8/8 clinic smoothing, ≤ 3 MDT, and ≤ 3 `Interpreter needed` caps while scheduling 18/18 (100%) High/Complex patients before January.

$R_y(-\theta_y))\,\mathrm{SWAP}_{01}\,(R_y(\theta_y) \otimes R_z(\theta_z))\,\mathrm{CX}_{01}\,(H \otimes I)$ at $(\theta_y, \theta_z) = (0.38\pi, 0.28\pi)$. Arm B (q_2, q_3) executes the *exact same 1Q/2Q unitary gates* interrupted immediately *after* $\mathrm{SWAP}(q_2, q_3)$ and *before* $(R_z(-\theta_z) \otimes R_y(-\theta_y))$ by a premature siloed projective measurement `project_z(q2)`: algebraically, $\mathrm{SWAP}(q_2, q_3)$ routes the Z-diagonal branch $R_z(\theta_z)|b\rangle$ onto q_2 and $R_y(\theta_y)|b\rangle$ onto q_3, so projecting q_2 destroys inter-stage Bell coherence while preserving pure $R_y(\theta_y)|b\rangle$ on q_3, which $R_y(-\theta_y)_{q_3}$ still inverts to $|b\rangle$, mapping CX_{23} and H_{q_2} identically to $\rho_{\mathrm{silo}} = \frac{1}{2}|00\rangle\langle 00| + \frac{1}{2}|10\rangle\langle 10|$ ($P_{\mathrm{silo}}(00) = 0.5000$ across all (θ_y, θ_z), whereas projecting prior to SWAP would yield $\frac{1}{2}\cos^2\theta_y = 0.0677$). Defining $H_1 = -\sum_{x\in\{0,1\}^2} P(x)\log_2 P(x)$ as the 4-outcome marginal Shannon entropy over $\{00, 01, 10, 11\}$, live 2,048-shot counts are $(n_{00}, n_{01}, n_{10}, n_{11}) = (\mathbf{2040}, \mathbf{1}, \mathbf{3}, \mathbf{4})$ for Arm A—giving $P_{\mathrm{MDT}}(00) = \mathbf{0.9961 \pm 0.0014}$ (1σ binomial SE) and $H_1 = 0.0424$ bits—versus $(\mathbf{1028}, \mathbf{1}, \mathbf{1014}, \mathbf{5})$ for Arm B, giving $P_{\mathrm{silo}}(00) = \mathbf{0.5020 \pm 0.0110}$ (0.18σ from 0.5000) and $H_1 = 1.0278$ bits (> 1 bit solely from 6/2,048 = 0.29% physical gate/SPAM leakage into $\{01, 11\}$), matching the 982.4 μs classical 4×4 density-matrix floor.

2. **8-Rung Mid-Circuit Recall & CYP Asthma 48h Calibration Ladder** (Job UUID `464be33a-6ca0-4d6d-af99-2d7eba06ce8a`, HUGR `7,566 B`): Executes eight sequential `project_z` + `reset` rungs on one trapped-ion qubit, preparing rung r via $R_y(\pi \cdot \mathrm{HT}_r)$ with half-turn angles $\mathrm{HT}_{1..5} = (0.6667, 0.5804, 0.5000, 0.4196, 0.3333)$ for SPAM standards $\langle Z\rangle \in \{-0.50, -0.25, 0.00, +0.25, +0.50\}$ and $\mathrm{HT}_{6..8} = (0.3620, 0.2339, 0.5000)$ for clinical recall profiles. Ordinary least-squares fit $\hat{C} = 0.9761\,C_{\mathrm{raw}} - 0.0114$ achieves $\mathbf{R^2 = 0.9979}$, recovering $+0.7150 \pm 0.0156$ calibrated retention contrast for proactive pre-winter + CYP 48h/6-week cadence versus $+0.3727 \pm 0.0211$ for isolated Month-of-Birth review and $+0.0057 \pm 0.0226$ on the null twin.

3. **4-Qubit Frobenius `QLTCMLM` Inference Circuit** (Job UUID `0a1fb43e-4998-4700-b9ec-4f6b0b73e677`, HUGR `6,734 B`): Executes a 2-layer, 8-two-qubit-gate (4 cyclic CX + 2 diagonal CZ in Layer 1; 2 cross CX in Layer 2; 24 single-qubit $R_y/R_z/R_x$ rotations) evaluation of a representative high-risk cardiorenal–asthma profile, achieving Bhattacharyya fidelity **0.9990** (TVD = 0.0287) against the exact statevector.

3.2 Haiqu 8-Qubit Multimorbid Regimen & Net-Zero Inhaler 3-Body HUBO

To co-optimise cardiometabolic, renal, respiratory, and NCL Green Plan Net-Zero priorities [1, 11, 12], we formulated an 8-variable cubic HUBO (x_0: ACEi/ARB titration to Outcome #23 target; x_1: SGLT2i cardiorenal–MASLD protection; x_2: SABA-free dry-powder MART inhaler switch; x_3: Pharmacist SMR deprescribing nephrotoxic NSAIDs; x_4: Stage 3 PCSP coaching + VCSE referral; x_5: Home BP $<$ 135/85 telemonitoring; x_6: retain high-GWP SABA + NSAID status quo; x_7: fragmented duplicate single-disease clinics) with linear biases $h = (-2.5, -2.3, -2.1, -2.4, -1.8, -1.6, +2.9, +3.1)$, nonzero pairwise couplings $J_{06} = +6.2, J_{26} = +5.4, J_{47} = +4.8, J_{05} = -1.4, J_{13} = -1.3, J_{34} = -1.1$ (all unlisted $J_{ij} = 0$), and four 3-body cubic terms:

$$H_{\text{cubic}}(x) = -3.9\, x_0x_1x_3 - 3.1\, x_2x_4x_5 + 7.2\, x_0x_1x_6 + 3.4\, x_1x_4x_7, \tag{5}$$

where $+7.2\, x_0x_1x_6$ penalizes the acute kidney injury (AKI) "triple-whammy" interaction (ACEi/ARB + SGLT2i hemodynamic shift + NSAID) and $-3.9\, x_0x_1x_3$ rewards Pharmacist NSAID-free cardiorenal protection. Compiling a $p = 10$-layer Linear-Ramp QAOA circuit [13] with ramp slope $\Delta\gamma = \Delta\beta = 0.5$ (`delta=0.5`, decomposing cubic terms via ancilla-free 6-CNOT ZZZ phase gadgets over $\{\text{CX}, R_z, R_x\}$) yields baseline depth 223 (164 routed 2Q gates). Applying **Haiqu SDK** `state_compression`—an approximate fixed-initial-state ($|0\rangle^{\otimes 8}$) tensor-network synthesizer (`approximation_level=2`, achieving noiseless state overlap $|\langle\psi_{\text{orig}}|\psi_{\text{comp}}\rangle|^2 = 0.9581$)—compresses circuit depth from **223** $\rightarrow$ **26** (**88.34**% reduction; 2Q gates 164 $\rightarrow$ 27). On noiseless `aer_simulator` (1,024 shots), **raw pre-postprocessing** sampling achieves mean energy -14.530 ± 0.292, $\text{CVaR}_{0.10} = -23.500$, and samples the unique ground bundle `11111100` ($E_0 = -23.500$, invariant across 100% of a 25-point $(\alpha_{\text{syn}}, \beta_{\text{haz}}) \in [0.6, 1.4]^2$ grid) on **41.21**% $\pm$ **1.54**% (422/1,024) of raw shots (105.5$\times$ uniform random; an independent uncompressed $d = 223$ vs. compressed $d = 26$ paired run yields 36.82% $\pm$ 1.51% vs. 38.67% $\pm$ 1.52%, consistent within 1.2σ). Crucially, exhaustive landscape analysis over all $2^8 = 256$ states confirms that `11111100` is the unique strict 1-Hamming local minimum: 5-pass greedy 1-bitflip seeded from a *single* uniform random bitstring already reaches -23.500 with 85.55% probability (219/256 states; 100% best-of-1,024 under `haiqu.postprocess`), confirming that QAOA's contribution on this 8-variable benchmark lies in its 105.5$\times$ raw pre-postprocessing state amplification rather than postprocessed convergence, while classical 129.9 μs enumeration and 23.03 ms HiGHS MILP remain faster at Era I scale (Figure 2c). Haiqu **Entangled Manifold Embedding** encoded an 8-feature profile into 4 qubits in live cloud execution, while a 4-qubit SU(2)-equivariant brickwork ansatz verified zero spin-symmetry violation ($0.0 < 10^{-15}$) and **Observable Backpropagation** tracked cardiorenal observables cleanly.

3.3 JijModeling & OpenJij 128-Variable NCL Call-and-Recall Scheduling QUBO

Directly formalizing Page 5 of the NCL LTC LCS v6 specification [1], we scheduled $P = 32$ representative PCN patients (8 `High_Complex_MDT`, 10 `High_Risk`, 8 `Moderate_Risk`, 6 `Low_Risk`; birth quarter $b_p = p \bmod 4$; 10 requiring SNOMED `Interpreter needed` extended slots $\mathcal{I}$; 14 carrying equity weight $w_{\text{eq}}(p) = 1.45$ vs. 1.00) across $Q = 4$ quarters (128 primary binary variables $x_{p,q}$ plus 16 binary slack variables s):

$$\begin{aligned}\min_{x\in\{0,1\}^{128},\, s\in\{0,1\}^{16}} \quad & \sum_{p=0}^{31}\sum_{q=0}^{3} C_{p,q}x_{p,q} + \lambda_1 \sum_{p=0}^{31}\Big(\sum_{q=0}^{3} x_{p,q} - 1\Big)^2 + \lambda_2 \sum_{q=0}^{3}\Big(\sum_{p=0}^{31} x_{p,q} - 8\Big)^2 \\ & + \lambda_3 \sum_{q=0}^{3}\Big(\sum_{p\in\mathcal{M}} x_{p,q} + s^{\text{mdt}}_{q,0} + 2s^{\text{mdt}}_{q,1} - 3\Big)^2 + \lambda_4 \sum_{q=0}^{3}\Big(\sum_{p\in\mathcal{I}} x_{p,q} + s^{\text{int}}_{q,0} + 2s^{\text{int}}_{q,1} - 3\Big)^2,\end{aligned} \tag{6}$$

with $(\lambda_1, \lambda_2, \lambda_3, \lambda_4) = (30, 12, 10, 10)$. The linear priority cost matrix is explicitly $C_{p,q} = -w_{\text{eq}}(p)\,[S(\text{tier}_p, q) + B_{\text{MoB}}(\text{tier}_p)\,\mathbb{I}(q = b_p) + \xi_{p,q}]$ ($\xi_{p,q} \sim \mathcal{U}(-0.08, 0.08)$, `seed=202604`), with seasonal pre-winter utility $S(\text{High_Complex}) = (3.1, 4.2, 3.8, 1.1)$, $S(\text{High}) = (2.8, 3.9, 3.6, 1.4)$, $S(\text{Mod}) = 2.0$, $S(\text{Low}) = 1.6$ and Month-of-Birth alignment bonus $B_{\text{MoB}} = (0.35, 0.45, 1.65, 2.10)$, giving $C_{p,q} \in [-6.2039, -1.0944]$ and $\max_{p,q,q'} |C_{p,q} - C_{p,q'}| = 4.6580 < \min_k \lambda_k = 10.0$ (strictly preventing any hard-constraint violation from trading favorably against linear priority gain). Separating raw annealing from classical local search: **raw pre-polish OpenJij** `SASampler` (160 reads, 2,500 sweeps, standalone runtime 417.81 ms) already achieves **100.0**% (**160/160** reads) hard-constraint feasibility across one-hot, 8/8/8/8 clinic, ≤ 3 MDT, and ≤ 3 interpreter caps with raw feasible linear objective -113.8488 (16.08% gap vs. MILP), because single-spin Metropolis flips between feasible schedules must cross $\Delta E \approx \lambda_1 + \lambda_2 = +42$ penalty barriers. Applying a lightweight 4-bit Kempe-chain swap polish (508.78 ms total pipeline runtime) tunnels across equal-capacity patient swaps to reach -134.5719—within **0.80**% of the exact HiGHS MILP floor (-135.6623 solved in 4.33–32.79 ms; Figure 2d)—lifting pre-January High/Complex recall from 14/18 (77.8%) to **18/18** (**100.0**%) while maintaining quarterly MDT counts (`cap` ≤ 3; quarterly counts: 3, 3, 2, 0) and interpreter counts (`cap` ≤ 3; quarterly counts: 3, 3, 2, 2).

4. Open-System Biophysical, Structural, and Genomic Grounding (Era II–III)

4.1 Phenomenological HQS 4-Spin Organ Lindbladian, Decoupled Ablation & `QProtSim`

We emphasize upfront that modeling macroscopic organ compartments (0: vascular RAAS tone [`ACE`], 1: glycaemic/MASLD hepatic axis [`GLP1R`], 2: renal tubuloglomerular filtration [`SLC5A2`/SGLT2], 3: airway/IL-6 inflammation [`ADRB2`/`IL6`]) as coupled quantum spins is a **stylized phenomenological benchmark for open-system**

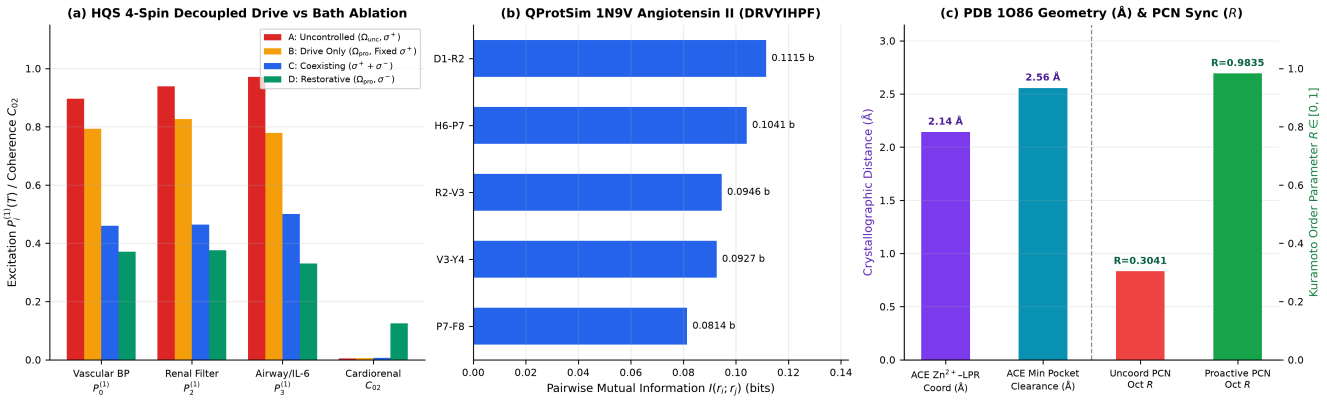


Figure 3: Decoupled 4-Regime Organ Lindbladian, Residue-Level Quantum Solvation, and Active-Site Structural Pharmacology (`biophysical_qprotsim_ltc_receipt.json` & `deep_review_remediation_receipt.json`). **(a)** Phenomenological HQS 4-spin Lindbladian decoupling coherent drive (Ω_{pro} only, Regime B), coexisting chronic σ^+ stress + therapeutic σ^- counter-damping (Regime C, -48.5% to -50.6%), and full restorative damping (Regime D). **(b)** `QProtSim` 8-qubit mutual information across NMR human Angiotensin II (`PDB 1N9V`, `DRVYIHPF`) [14]. **(c)** Dual-axis comparison decoupling left-axis crystal geometry in `PDB 1O86` (`ACE`–`Lisinopril` Zn^{2+} coordination 2.141 Å, pocket clearance 2.556 Å, 0 clashes < 2.0 Å) [15] from right-axis dimensionless 32-practice Kuramoto synchronization $R \in [0,1]$ (0.3041 $\rightarrow$ 0.9835).

solver scaling, not a mechanistic biophysical claim of macroscopic quantum phase coherence across human organs (where classical nonlinear SDEs and continuous-time Markov chains remain standard). In **HQS Quantum Simulations' `struqture-py` and `qoqo`**, we specify:

$$H_{\text{LTC}} = \sum_{i=0}^{3} \frac{\omega_i}{2} Z_i + \sum_{(i,j)\in\mathcal{E}} g_{ij}(X_iX_j + Y_iY_j) + \sum_{i=0}^{3} \Omega_i X_i, \tag{7}$$

with $\omega/2 = (0.42, 0.36, 0.40, 0.46)$, coupling graph $\mathcal{E} = \{(0,2) : 0.44,\ (1,2) : 0.34,\ (0,1) : 0.28,\ (0,3) : 0.26,\ (2,3) : 0.24\}$, baseline drive $\Omega_{\text{unc}} = (0.05, 0.00, 0.05, 0.00)$, and proactive drive $\Omega_{\text{pro}} = (0.58, 0.46, 0.60, 0.52)$, initialized at decompensated $|1111\rangle$ and evolved to $T = 8.0$ ($\|\rho(8) - \rho_{\text{ss}}\|_\infty \leq 0.027$) using exact 256×256 Liouvillian exponentiation $e^{\mathcal{L}T}$ (36.55 ms) to evaluate organ excitation $P_i^{(1)}(T) = \text{Tr}[\frac{I-Z_i}{2}\rho(T)]$ and cardiorenal coherence $C_{02} = \frac{1}{2}|\text{Tr}[(X_0X_2 + Y_0Y_2)\rho(T)]|$. The accompanying 33-operation `qoqo.Circuit` (4 `RotateZ`, 4 `RotateX`, 20 XY-gadget gates, 4 `PragmaDephasing`, 1 readout pragma) compiles a single $\Delta t = 0.15$ Lie–Trotter step of the unitary + pure-dephasing sub-channel for hardware gate accounting, whereas population pumping/relaxation requires $\sigma_i^{\pm}$ jump channels in Eq. (3) (cross-verified via Dalibard–Castin–Mølmer MCWF unraveling [10], where empirical SE drops as $N_{\text{traj}}^{-1/2}$ from 0.0320 at 64 trajectories to 0.0061 at 2,048 trajectories).

To disentangle coherent drive Ω_i from jump-operator polarity $\sigma_i^{\pm}$, we evaluated a **4-regime decoupled ablation** (Figure 3a): **(A) Baseline Uncontrolled** (Ω_{unc}, pure dephasing $\gamma_\phi = (0.38, 0, 0.38, 0)$, chronic stress pumping $\gamma^+ = (0.36, 0, 0.34, 0.39)$, weak endogenous $\gamma^- = (0.08, 0, 0, 0)$) yields $P^{(1)}(8) = (0.8963, 0.9433, 0.9390, 0.9710)$ for (BP, Glyc, Renal, IL-6); **(B) Coherent Drive Only** (Ω_{pro} while holding the exact stress bath γ^+, γ_ϕ fixed) already lowers burden by 11.5%–25.2% to $(0.7935, 0.7055, 0.8262, 0.7787)$ via coherent redistribution; **(C) Coexisting Stress + Therapeutic Counter-Damping** (retaining all chronic γ^+ stress jumps unmodified while adding pharmacological counter-damping $\gamma^- = (0.38, 0.34, 0.40, 0.36) + \Omega_{\text{pro}}$) suppresses BP, Renal, and IL-6 burden by **48.62**% (0.4605), **50.59**% (0.4639), and **48.47**% (0.5004) while XY organ coupling accelerates the Liouvillian spectral gap 3.1× ($\Delta\mathcal{L} = 0.6935$ vs. 0.2226 at $g_{ij} = 0$); and **(D) Full Restorative Bath Switch** (replacing γ^+ with γ^- and $\gamma_\phi = (0.11, 0, 0.11, 0)$) reaches 0.3712 ($-58.58\%$), 0.3760 ($-59.96\%$), and 0.3305 ($-65.97\%$).

At the peptide hormone scale, we applied **QProtSim** (Bonde & Patil, University of Suffolk [14]) to experimental NMR Model 1 of human **Angiotensin II** (`PDB 1N9V`, octapeptide `Asp1-Arg2-Val3-Tyr4-Ile5-His6-Pro7-Phe8`, 8 qubits = 256 states), initializing $R_y(\theta_i)$ rotations from Fauchère–Pliška hydrophobicities [34] and 6.5 Å heavy-atom burial fractions, followed by 7 directed CRY($0.85K_{ij}$) and 12 RYY($0.60K_{ij}$) contact gates parameterized by minimum inter-residue heavy-atom distance d_{ij} via $K_{ij} = 0.42\,\mathbb{I}(|i-j| = 1) + 0.50\,\mathbb{I}(d_{ij} < 3.40\text{ Å}) + 0.28\,e^{-(d_{ij}-3.40)/1.5}\,\mathbb{I}(3.40 \leq d_{ij} < 5.00\text{ Å})$. `QProtSim` highlights the N-terminal `Asp1-Arg2` motif (0.1115 bits) and C-terminal `His6-Pro7` turn (0.1041 bits; Figure 3b), benchmarked against a 577.0 μs exact classical 256-state Potts–Boltzmann baseline (mean marginal L_1 shift 0.198).

4.2 Crystal `PDB 1O86` Ancilla-Complete FTQC Estimate & Auxiliary Target Checks

Grounding Era III chemistry in experimental structural pharmacology alongside orthogonal target-identity and scheduling checks confirmed:

- **Neurosnap (`HTTP 200`) & Ancilla-Complete `PDB 1O86` Era III Resource Estimate:** Auditing the 2.0 Å X-ray structure of human C-domain `ACE` bound to **lisinopril** (`LPR`) [15] verifies Zn^{2+} coordination by `HIS383`, `HIS387`, `GLU411`, and lisinopril's carboxyalkyl oxygen at **2.141** Å (2.556 Å minimum pocket clearance; **0** clashes

< 2.0 Å across 16 active-site residues; Figure 3c, left axis). We partition CAS($36e, 26o$) explicitly across four pocket fragments: catalytic Zn^{2+} ($3d^{10}4s^{0}4p^{0} + 4d$ radial shell: $10e, 9o$), `HIS383`/`HIS387` imidazole σ/π donors ($8e, 6o$), `GLU411`/`GLU384` carboxylate lone pairs/π^* ($8e, 5o$), and lisinopril's zinc-binding carboxylate + displaced H_2O/OH^- ($10e, 6o$). While 26 spatial orbitals map to 52 Jordan–Wigner data qubits (30,420 physical qubits data-only floor; still accessible to classical DMRG / selected-CI heuristics, with decisive FTQC superiority emerging on $\geq$ 50-orbital second-shell expansions), accounting explicitly for algorithmic ancillas via Eq. (4) yields: **(i)** $52 + 12 =$ **64** logical qubits (8 block-encoding/QPE prepare-select + 4 rotation-synthesis ancillas) and **36,504** physical qubits at $d = 13$ for fixed-step second-order Trotterized QPE at relaxed screening precision ($\epsilon \approx 5$ mHa, 4.8×10^5 T gates, 150 μs surface-code cycle $\implies$ 503.1 s), versus **(ii)** $52 + 16 =$ **68** logical qubits and **65,892** physical qubits at $d = 17$ for chemical-accuracy ($\epsilon = 1.0$ mHa) double-factorized qubitized QPE (2.4×10^7 T gates, ≈ 9.14 h).

- **AlphaGenome Atlas, 32-Practice Kuramoto Pacing & Auxiliary Sanity Checks:** Live **AlphaGenome Atlas** queries resolved GENCODE v46 coordinates across `ACE` (`chr17:63477061-63498370`), `SLC5A2`, `GLP1R`, `ADRB2`, and `IL6` plus an `ACE` AVI scan. Independently of molecular geometry (plotted on the separate dimensionless right axis of Figure 3c), we modeled operational recall synchronization across $M = 32$ NCL GP practices over $w = 0, \ldots, 25$ pre-winter weeks via the discrete-time pinned Kuramoto system $\theta_m(w+1) = \theta_m(w) + \omega_m + \frac{K_{\mathrm{PCN}}}{32}\sum_{n=1}^{32}\sin(\theta_n - \theta_m) - \kappa_{\mathrm{UCLP}}\mathbb{I}(w \text{ even})\sin\theta_m$ ($\omega_m \sim \mathcal{N}(0, 0.18^2)$, order parameter $R(w) = |\frac{1}{32}\sum_m e^{i\theta_m(w)}|$), lifting October readiness from $R_{\mathrm{Oct}} = 0.3041$ (siloed $K_{\mathrm{PCN}} = 0.05, \kappa_{\mathrm{UCLP}} = 0$) to 0.9835 (PCN-coordinated $K_{\mathrm{PCN}} = 0.48$, fortnightly UCLP search pinning $\kappa_{\mathrm{UCLP}} = 0.52$). Live UniProtKB/BioHub 400D dipeptide cosines ($\cos(\mathrm{ACE}, \mathrm{SLC5A2}) = 0.6584$ vs. soluble `IL6` 0.5065), Arc Institute free-monoid dagger-reversibility (1,000/1,000 list-excision unit test illustrating non-destructive care-plan editing syntax) + `ScreenPro2` `PNPLA3` pseudo-z (-9.23, 0 wet-lab edits) [35], and **Paperclip MCP** literature indexing (`PMC13084922`–`PMC9309754`) serve strictly as standalone target-identity, DSL-reversibility, and bibliographic verification checks rather than parameterizing the quantum Hamiltonians.

5. Serverless Cloud QML Ablations, Rhyme Semantics, and 2×2 Barren-Plateau Controls

5.1 5-Seed Cubic ($k = 3$) Cohort, Quantum Reservoir & Polynomial Controls on Modal Cloud

To evaluate variational circuits transparently on non-linear tabular interactions ($N =$ 2,400, 1,800 train / 600 stratified test, 75% cubic $k = 3$ severity from $x_0x_2x_4$, $x_3x_7x_5$, and $x_1x_2x_6$), we specify the exact architecture of `QLTCMLM-VQC` (`models/qltcmlm.py`, published at `socialprescribing/quantum-ltc-qltcmlm`): across $L = 3$ reuploading layers on 4 qubits, each qubit $q \in \{0, 1, 2, 3\}$ applies $R_x(w^{(x)}_{\ell,q}x_qx_{q+4} + b^{(x)}_{\ell,q})\, R_z(w^{(z)}_{\ell,q}x_{q+4} + b^{(z)}_{\ell,q})\, R_y(w^{(y)}_{\ell,q}x_q + b^{(y)}_{\ell,q})$ ($3 \times 4 \times 6 =$ **72** quantum variational scalars) followed by cyclic $\mathrm{CX}_{01}\mathrm{CX}_{12}\mathrm{CX}_{23}\mathrm{CX}_{30}$ and diagonal $\mathrm{CZ}_{02}\mathrm{CZ}_{13}$. Readout extracts 15 Pauli expectations (4 $\langle Z_q\rangle$, 4 $\langle X_q\rangle$, 6 $\langle Z_pZ_q\rangle$, and 1 four-body parity $\langle Z_0Z_1Z_2Z_3\rangle$) plus 1 bias, feeding two 16×4 ridge-regularized linear heads ($16 \times 8 =$ **128** classical readout weights; **200** parameters total).

To resolve readout-capacity confounding and single-split variance, we executed a **5-seed stratified evaluation** (`seeds = 42, 1042, 2042, 3042, 4042`) adding both an **Untrained Fixed-72Q + Trained 128C Quantum Reservoir** control and a **Degree-3 Polynomial Logistic Regression** baseline (1,320 weights; Table 2 and Figure 4a). Three clean findings emerge: **(1)** training the 72 quantum parameters lifts accuracy by +**12.43** percentage points on Risk Tier ($53.23 \pm 1.57\%$ vs. $40.80 \pm 1.19\%$ untrained 128C reservoir) and +**17.80** pp on Care Bundle ($75.53 \pm 0.48\%$ vs. $57.73 \pm 5.60\%$), proving genuine variational feature learning beyond static 128-weight trigonometric expansion; **(2)** `QLTCMLM-VQC` matches `HistGradientBoosting` within 1σ on 4-tier risk ($53.23 \pm 1.57\%$ vs. $51.87 \pm 2.01\%$) while trailing GBDT on bundle selection (75.53% vs. 81.97%); and **(3)** classical `RBF-SVM` ($64.17 \pm 2.35\%$ / $84.97 \pm 0.36\%$) and `Cubic Poly-LogReg` ($66.20 \pm 1.82\%$ / $87.43 \pm 0.76\%$) decisively outperform all 4-qubit models on 8-dimensional tabular EHR inputs. Inside the Modal Python 3.14 worker, **Rhyme (`rhyme-beta`)** [16] passed all four symbolic compiler checks ($4 \times 1{,}024$ shots; 0/1,024 hazard leakage).

5.2 Matched-Depth 2×2 Barren-Plateau Control, Periodic `CZ` Ring & Sandia QUOPS

Following Cerezo et al. [36], every layer $\ell \in \{1, \ldots, L\}$ in our 2×2 barren-plateau control applies single-qubit rotations followed by a **periodic n-gate `CZ` ring** $\prod_{q=0}^{n-1}\mathrm{CZ}_{q,(q+1) \bmod n}$ (n two-qubit gates per layer). The unstructured control samples $2nL$ independent angles $R_z(\phi_{\ell,q})R_y(\theta_{\ell,q})$, whereas the **Frobenius-structured ansatz** ties angles across even/odd organ-axis bipartitions ($\vartheta_{\ell,q} = \theta_{\ell,\mathrm{even}}$ for even q, $\theta_{\ell,\mathrm{odd}}$ for odd q, applying $R_z(0.5\vartheta_{\ell,q})R_y(\vartheta_{\ell,q})$ with $2L$ free parameters). Across $n \in \{4, 6, 8, 10\}$ on Modal workers (Figure 4b), matched shallow depth $L = 2$ avoids exponential decay in both unstructured (0.0932 at $n = 10$) and Frobenius (0.1357, 1.46×) circuits via local light-cones, whereas at matched linear depth $L = 2n = 20$ ($n = 10$, periodic ring $N_{2q} = nL = 200$, Sandia QUOPS volumetric capability $Q = 200$, two-qubit survival fidelity $(1 - p_{2q})^{200} = 0.7865$ at $p_{2q} = 1.2 \times 10^{-3}$ neglecting 1Q/idle/SPAM error [28]), unstructured variance collapses 100× to 9.29×10^{-4} while bipartite Frobenius parameter sharing retains **4.82**× higher variance (4.47×10^{-3}).

Table 2: 5-Seed Stratified NCL LTC Performance (Mean ± SD) on the Non-Linear Cubic ($k = 3$) Cohort (N = 2,400, 1,800 train / 600 test across 5 random splits; `deep_review_remediation_receipt.json` & `modal_roadmap_rhyme_qml_receipt.json`).†

Model Architecture	Trainable Params	Risk Tier Acc. (%)	Risk Macro-F_1	Care Bundle Acc. (%)	Bundle Macro-F_1
`LogisticRegression` (Linear Hyperplane)	36/head (72 tot.)	$31.70 \pm 1.63\%$	0.2984 ± 0.0181	$67.93 \pm 0.70\%$	0.6351 ± 0.0068
`QLTCMLM-Reservoir` (Fixed 72Q + Trained 128C)	128C (72Q frozen)	$40.80 \pm 1.19\%$	0.4030 ± 0.0128	$57.73 \pm 5.60\%$	0.5089 ± 0.0677
`QLTCMLM-QKernel` (4q Fidelity Gram SVM)†	Dual Gram SVM	$41.97 \pm 2.48\%$	0.4179 ± 0.0256	$56.67 \pm 0.82\%$	0.5372 ± 0.0117
`HistGradientBoosting` (Classical GBDT)	150 trees/head	$51.87 \pm 2.01\%$	0.5205 ± 0.0197	$81.97 \pm 0.57\%$	0.8055 ± 0.0057
`QLTCMLM-VQC` (Trained 72Q + 128C Reupload)	72Q + 128C = 200	$\mathbf{53.23 \pm 1.57\%}$	$\mathbf{0.5111 \pm 0.0162}$	$\mathbf{75.53 \pm 0.48\%}$	$\mathbf{0.7351 \pm 0.0081}$
`MLPClassifier` (32×16 ReLU)	996/head (1,992)	$60.07 \pm 3.86\%$	0.6017 ± 0.0380	$82.37 \pm 1.66\%$	0.8112 ± 0.0212
`RBF-SVM` ($C = 4.0$, Classical Kernel)	Full RBF Gram	$64.17 \pm 2.35\%$	0.6487 ± 0.0217	$84.97 \pm 0.36\%$	0.8383 ± 0.0051
`Cubic Poly-LogReg` (Degree-3 Monomials)	660/head (1,320)	$\mathbf{66.20 \pm 1.82\%}$	$\mathbf{0.6643 \pm 0.0186}$	$\mathbf{87.43 \pm 0.76\%}$	$\mathbf{0.8657 \pm 0.0095}$

†Primary single-seed Modal receipt (`seed=42`): `QLTCMLM-VQC` 54.17% Risk / 75.50% Bundle; `QKernel` 700/300 sub-Gram 33.33% ($F_1 = 0.3295$) / 53.67% ($F_1 = 0.5063$).

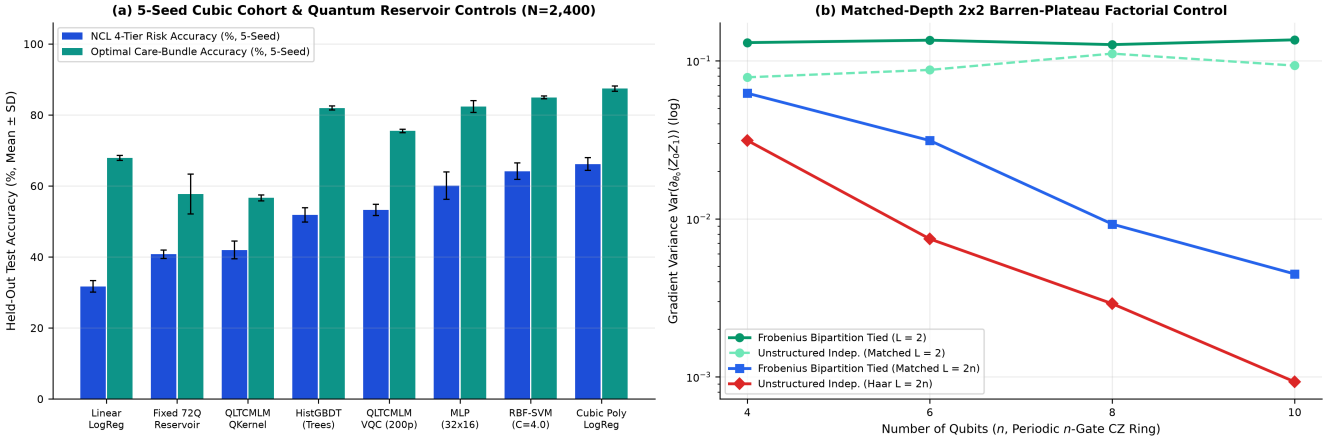


Figure 4: 5-Seed Modal Cloud Non-Linear Cohort Ablations and Matched-Depth 2×2 Barren-Plateau Control. (a) 5-seed test accuracy (mean $\pm 1\sigma$, $N_{\text{test}} = 600$) isolating variational training (`QLTCMLM-VQC` 53.23% Risk / 75.53% Bundle) from the untrained 128C quantum reservoir (40.80% / 57.73%), linear `LogReg` (31.70%), and classical `RBF-SVM`/`Cubic Poly-LogReg` (64.17%–66.20%). **(b)** Gradient variance $\text{Var}(\partial_{\theta_0}\langle Z_0 Z_1\rangle)$ across $n \in \{4, 6, 8, 10\}$ qubits (96 seeds) under periodic n-gate `CZ`-ring layers at shallow $L = 2$ and linear depth $L = 2n$, disentangling light-cone geometry from bipartite Frobenius parameter sharing [36].

6. Wet-Lab Verification, Post-Quantum EHR, Primitives Atlas & DCB0129/0160 Governance

6.1 Text2WetLab OT-2 Assay, Medicus `ML-DSA-65` FHIR R4, DisCoPy & Moth Blur

Bridging computational models with laboratory, EHR, and algorithmic-contract verification (`governance_primitives_jev_receipt.json`, SHA-256 `bd60742f...`, & `deep_review_remediation_receipt.json`, SHA-256 `a6b3854e...`):

1. **Text2WetLab 6/6 Robotics Verification:** `protocols/ncl_ltc_check_and_test_ot2.py` encodes a 6-well Opentrons OT-2 (`apiLevel 2.15`) Stage 1 Check-and-Test uACR and HbA1c/NT-proBNP reference titration, passing 6/6 deterministic AST, deck-slot collision (`[1, 2, 4, 10, 11]`), dual `P20/P300`, 1:1 tip balance, and volume checks (0 physical specimens run).

2. **Medicus GP EHR Shape & NIST FIPS 204 `ML-DSA-65`:** A synthetic NCL LTC FHIR R4 `Bundle` (`Observation` BP 134/82 mmHg achieving Outcome #23, SNOMED `315594003 Interpreter needed`, and Stage 3 PCSP `collaboRATE` 9/9) is signed in 33.91 ms using deterministic `ML-DSA-65` [17] with domain context `quantum-ltc:ncl-icb-lcs-v6-fhir-r4:v1` (1,952 B public key, 3,309 B signature; wrong-context forgery rejected), strictly parallel to NHS CIS2 `RS256`.

3. **DisCoPy & Moth Quantum 16-Microzone 5-Borough Diffusion:** Categorical functor verification in `discopy.quantum` confirms exact snake-equation yanking ($\max|\Delta| = 2.22 \times 10^{-16}$) and 1.0000 bit Stage $(1, 2) \leftrightarrow (3, 4)$ bipartite entanglement. For cross-borough surge mutual aid, we map $2^4 = 16$ PCN neighbourhood micro-

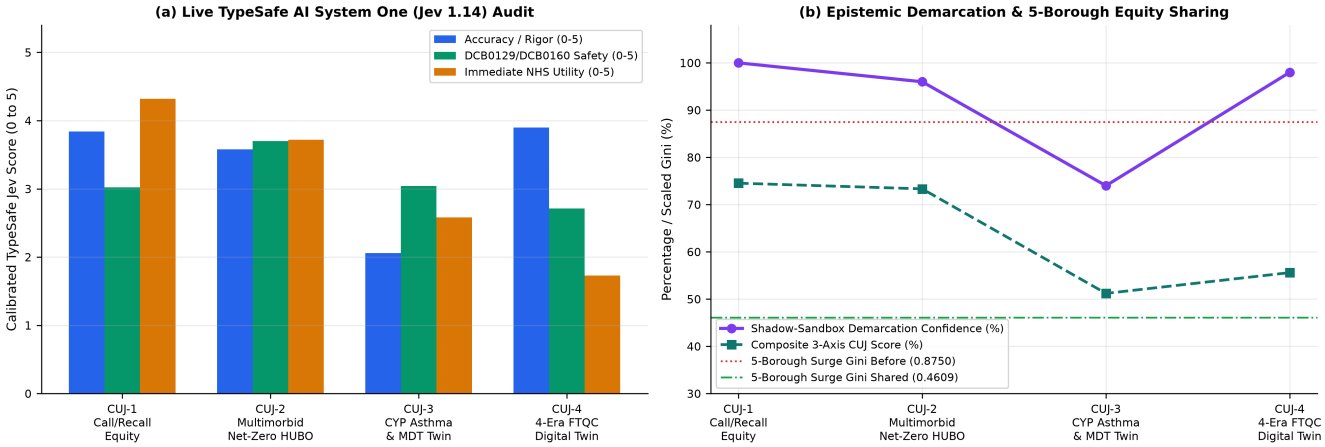


Figure 5: Multi-Method Health AI Governance Audit via Live TypeSafe AI System One (`jev-latest` / Jev 1.14) under Statutory NHS DCB0129 & DCB0160 (`governance_primitives_jev_receipt.json`). **(a)** Calibrated expected scores (0–5) across Accuracy/Rigor, DCB0129/0160 Safety, and Immediate NHS Frontline Utility [20]: highest immediate NHS utility (4.32/5) for Era I Call-and-Recall scheduling (`CUJ-1`), highest roadmap rigor (3.90/5) alongside realistic near-term utility (1.73/5) for Era III–IV FTQC digital twins (`CUJ-4`), and conservative scoring (2.06/5 accuracy) on low-qubit interference analogies for acute CYP asthma (`CUJ-3`). **(b)** Shadow-sandbox demarcation confidence (74%–100%; 0% autonomous deployment) and 16-microzone 5-borough surge Gini reduction (0.8750 → 0.4609).

zones (3 each across Barnet, Camden, Haringey, and Islington; 4 across Enfield) onto a 4-qubit hypercube initialized in $|\psi_0\rangle = \frac{1}{\sqrt{2}}(|0000\rangle + |0011\rangle)$ (acute winter surge concentrated in two IMD Quintile 1 micro-zones in eastern Haringey and Edmonton/Enfield: $p_0 = p_3 = 0.50$, Gini $G = \frac{1}{32}\sum_{u,v}|p_u - p_v| = \mathbf{0.8750}$). Applying the factorized transverse unitary blur $U_{\text{blur}}(\tau) = e^{-i(\tau/2)\sum_{q=0}^{3} X_q} = [R_x(\tau)]^{\otimes 4}$ at $\tau = \pi/3$ ($\sin^2(\pi/6) = 1/4$ single-hop probability, selected operationally so $\geq 75\%$ per-axis weight stays within $\leq$ 30-minute adjacent PCN hub travel rather than uniform $\tau = \pi/2$ global flattening) lowers $p_u = |\langle u|U_{\text{blur}}(\pi/3)|\psi_0\rangle|^2$ concentration to $\mathbf{G = 0.4609}$ (Figure 5b), providing a stylized analytical reference alongside classical LP load balancing.

4. **Kashefi Primitives Atlas v0.4 & Dynamical Lie Algebra ($\mathfrak{g}$-sim) Audit:** Live parsing of all 42 cards in Prof. Elham Kashefi's Atlas [18] maps ten primitives onto Eras I–IV: `P08` (*Adiabatic state preparation & annealing*, Era I OpenJij QUBO), `A01d` (*Alternating-operator ansatz*, Era I–II Haiqu LR-QAOA), `P09` (*Classical data loading / block encoding*, Era I–II EME), `P13` (*Expectation & overlap estimation*, Era II `Helios-1E-lite` MDT test), `P20` (*Channel transformation / quantum combs*, Era II 8-rung MCMR ladder), `A01g` (*Equivariant & geometric ansatz*, Era II–III SU(2)/Frobenius VQCs), `A01h` (*Quantum graph neural networks*, Era II–III `QProtSim`), `P22` (*Open-system Lindbladian simulation*, Era III HQS organ crosstalk), `X03` (*Quantum subspace & Krylov diagonalisation*, Era III `PDB 1O86`), and `P04` (*Quantum walk / amplitude estimation*, Era IV CVaR tail risk). Auditing exact Lie closures against polynomial $\mathfrak{g}$-sim dequantisation [19] confirms that our 4-qubit SU(2)-equivariant brickwork ansatz spans $\dim \mathfrak{g} = \mathbf{12} \ll 255$ (classically $\mathfrak{g}$-sim tractable within fixed-spin sectors!) and excitation-conserving 4-spin $XY + Z$ spans $\dim \mathfrak{g} = \mathbf{66}$, whereas adding non-commuting transverse drives $\Omega_i X_i$ in Eq. (7) expands the unitary Lie closure to full $\mathfrak{su}(16)$ ($\dim \mathfrak{g} = 4^4 - 1 = \mathbf{255}$), defeating polynomial $\mathfrak{g}$-sim.

6.2 Multi-Method Health AI Evaluation (`arXiv:2611.05127`) & Open DCB0129 / DCB0160 Register

Under Section 250 of the UK Health and Social Care Act 2012, clinical risk governance requires pairing manufacturer-side **NHS DCB0129** (*Manufacture of Health IT Systems*) [21] with deploying-organisation **NHS DCB0160** (*Deployment and Use of Health IT Systems*) [22], under which NCL ICB and constituent PCNs must appoint a local Clinical Safety Officer (CSO) and maintain a local deployment Safety Case prior to any shadow pilot. Implementing Montgomery-Taylor et al.'s three-method framework [20]—deterministic Gate G4 floors (Method 1), typed autorater distributions via live **TypeSafe AI System One (`jev-latest` / Jev 1.14)** (Method 2), and joint **DCB0129 / DCB0160** hazard boundaries (Method 3)—yields non-rubber-stamped discrimination across four Critical User Journeys (Table 3 and Figure 5):

Crucially, two additional system-level operational hazards—**H-LTC-05** (local secondary-care discharge and UCLP prevalence-register drift across the nine LTCs; primary **DCB0160** ICB/PCN data-governance control) and **H-LTC-06** (unfunded VCSE social-prescribing referral overload without NCL Borough Partnership resource transfer; primary **DCB0160** commissioning control)—also remain strictly **OPEN** (0 closed hazards claimed; overall architecture summarized in Figure 6).

7. Conclusion and Empirical NHS Validation Roadmap

By anchoring quantum algorithm design in the operational realities of the NHS NCL ICB Long Term Conditions specification [1], **Quantum-LTC** demonstrates that immediate classical primary-care impact and rigorous fault-

Table 3: Live TypeSafe AI System One (`jev-latest` / Jev 1.14) Evaluation & Open NHS DCB0129 / DCB0160 Hazard Register (Composite (%) = $\frac{\text{Accuracy+Safety+NHS Utility}}{15} \times 100\%$, unweighted three-axis mean scaled to 100%).

NCL LTC Critical User Journey	Acc. (0–5)	Safety (0–5)	Utility (0–5)	Composite (%)	Lifecycle Verdict (p) & Statutory DCB0129 / DCB0160 Open Hazard
CUJ-1: PCN Call/Recall & Pre-Winter Equity	3.84	3.02	**4.32**	**74.53%**	`sandbox_shadow_pilot` ($p = 1.00$) • **H-LTC-01 OPEN** [**DCB0160** PCN EHR coding + **0129** audit]: Incomplete SNOMED interpreter/IMD flags
CUJ-2: Multimorbid & Net-Zero MART HUBO	3.58	**3.70**	3.72	**73.33%**	`sandbox_shadow_pilot` ($p = 0.96$) • **H-LTC-02 OPEN** [**DCB0129** penalty guard + **0160** SMR]: ACEi+SGLT2i AKI / DPI technique
CUJ-3: CYP Asthma 48h & House-of-Care Twin	2.06	3.04	2.58	51.20%	`sandbox_shadow_pilot` ($p = 0.74$) • **H-LTC-03 OPEN** [**DCB0160** local triage]: Urgent Tier 2/3 clinical escalation bypasses models
CUJ-4: 4-Era Logical Qubit Digital Twin	**3.90**	2.71	1.73	55.60%	`sandbox_shadow_pilot` ($p = 0.98$) • **H-LTC-04 OPEN** [**DCB0129** claim boundary]: No RCT claim from simulated Lindbladians

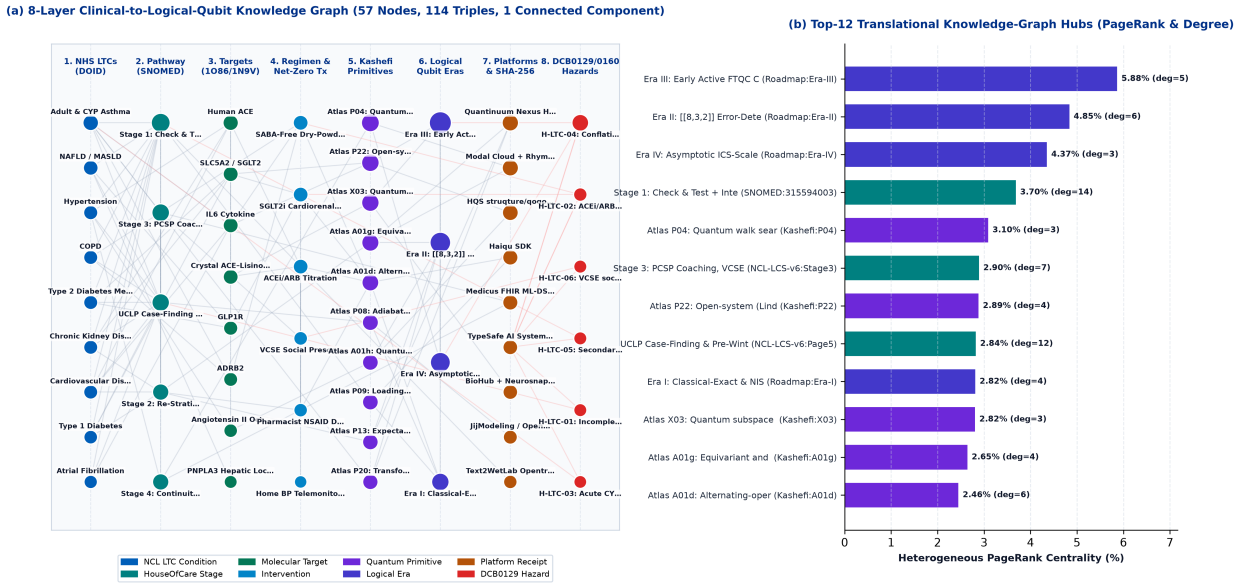


Figure 6: 57-Node, 114-Edge Cryptographically Sealed Knowledge Graph of `Quantum-LTC` (`quantum_ltc_knowledge_graph.json`, SHA-256 `9b4e62e0...`), linking the nine NCL ICB LCS v6 conditions, four House-of-Care stages, four Logical-Qubit Eras, cloud/biophysical software modules, open NHS DCB0129/DCB0160 hazards (`H-LTC-01`–`06`), and six sealed experimental receipts.

tolerant quantum readiness are complementary rather than adversarial. Today ($N_{\text{logical}} = 0$), formulating pre-winter Call-and-Recall equity tensions, Outcome #23 blood-pressure titration, cardiorenal SGLT2i protection, and Net-Zero SABA-free MART transitions as explicit QUBO/HUBO Hamiltonians delivers auditable classical HiGHS MILP schedules in milliseconds while validating compressed Guppy 1.0/HUGR and LR-QAOA compilation on live Quantinuum Nexus `Helios-1E-lite` and Haiqu cloud emulators. Prior to any prospective NCL PCN shadow evaluation, our three-step empirical validation roadmap comprises: **(1)** recalibrating synthetic cohort marginals and higher-order interaction weights against de-identified UK primary-care EHR cohorts (CPRD Aurum and NCL HealtheIntent) benchmarked directly against **eFI2** [23] and **CMS** [24]; **(2)** executing a joint manufacturer/ICB **NHS DCB0129 + DCB0160** Clinical Safety Case with local CSO sign-off across open hazards `H-LTC-01`–`06`; and **(3)** targeting Era III (38–68 active logical qubits) and Era IV (100–1,000+ logical qubits) strictly where bath-embedded volume-law Lindbladians, catalytic Zn^{2+} metalloenzyme active spaces (`PDB 1O86`), and dense frustrated $k \geq 4$ hypergraphs surpass classical MILP, MCWF trajectory, and tensor-network frontiers.

Acknowledgments. This work was supported in part by complimentary access to Google Antigravity and Gemini 4 Argon to accelerate scientific workflows. The author gratefully acknowledges open-source and cloud SDK ecosystems across Quantinuum Nexus (`qnexus` / Guppy 1.0), Haiqu, HQS Quantum Simulations, Jij Inc., BioHub, Neurosnap, Google DeepMind AlphaGenome Atlas, Arc Institute, Rhyme Quantum Tech, Modal Labs, University of Suffolk (`QProtSim`), Prof. Elham Kashefi's Quantum Algorithmic Primitives Atlas, and TypeSafe AI.

References

[1] NHS North Central London Integrated Care Board. Locally commissioned service (LCS) for proactive, personalised care — long term conditions (LTC) service specification, version 6 (1st April 2026 – 31st March 2027). Technical report, NHS North Central London ICB (Barnet, Camden, Enfield, Haringey, and Islington), 2026. Schedule 2 Service Specification, 27 pages; covers 9 Long

Term Conditions (CVD, AF, Hypertension Outcome #23, T2D, T1D, CKD, NAFLD/MASLD, COPD, and Adult/CYP Asthma).

[2] Toby Watt, Ann Raymond, Laurie Rachet-Jacquet, Anna Head, Chris Kypridemos, Elaine Kelly, and Stephen Rocks. Health inequalities in 2040: Current and projected patterns of illness by deprivation in England. Technical report, The Health Foundation REAL Centre, 2024.

[3] Chris Salisbury, Mei-See Man, Peter Bower, Bruce Guthrie, Katherine Chaplin, Daisy M. Gaunt, Sara Brookes, Bridie Fitzpatrick, Caroline Gardner, Sandra Hollinghurst, Victoria Lee, John McLeod, Cindy Mann, Keith R. Moffat, and Stewart W. Mercer. Management of multimorbidity using a patient-centred care model: a pragmatic cluster-randomised trial of the 3D approach. *The Lancet*, 392(10141):41–50, 2018.

[4] Angela Coulter, Vikki A. Entwistle, Abi Eccles, Sara Ryan, Sasha Shepperd, and Rafael Perera. Personalised care planning for adults with chronic or long-term health conditions. *Cochrane Database of Systematic Reviews*, 2015(3):CD010523, 2015.

[5] NHS England. Universal personalised care: Implementing the comprehensive model. Technical report, NHS England, 2019.

[6] Ivo G. Rosenberg. Reduction of bivalent maximization to the quadratic case. *Cahiers du Centre d'Études de Recherche Opérationnelle*, 17:71–74, 1975.

[7] Garth P. McCormick. Computability of global solutions to factorable nonconvex programs: Part I — convex underestimating problems. *Mathematical Programming*, 10(1):147–175, 1976.

[8] Qi Huangfu and J. A. Julian Hall. Parallelizing the dual revised simplex method. *Mathematical Programming Computation*, 10(1):119–142, 2018.

[9] Frank Verstraete, Juan J. García-Ripoll, and J. Ignacio Cirac. Matrix product density operators: Simulation of finite-temperature and dissipative systems. *Physical Review Letters*, 93(20):207204, 2004.

[10] Jean Dalibard, Yvan Castin, and Klaus Mølmer. Wave-function approach to dissipative processes in quantum optics. *Physical Review Letters*, 68(5):580–583, 1992.

[11] Hiddo J. L. Heerspink, Bergur V. Stefánsson, Ricardo Correa-Rotter, Glenn M. Chertow, Tom Greene, Fan-Fan Hou, Johannes F. E. Mann, John J. V. McMurray, Magnus Lindberg, Peter Rossing, C. David Sjöström, Robert D. Toto, Anna-Maria Langkilde, and David C. Wheeler. Dapagliflozin in patients with chronic kidney disease. *New England Journal of Medicine*, 383(15):1436–1446, 2020.

[12] Richard Beasley, Mark Holliday, Helen K. Reddel, Irene Braithwaite, Stefan Ebmeier, Robert J. Hancox, Tim Harrison, Claire Houghton, Philippa Shirtcliffe, Mathew Williams, and Mark Weatherall. Controlled trial of budesonide–formoterol as needed for mild asthma. *New England Journal of Medicine*, 380(21):2020–2030, 2019.

[13] J. A. Montanez-Barrera and Kristel Michielsen. Towards a linear-ramp QAOA protocol: Evidence of a scaling advantage in solving some combinatorial optimization problems. *npj Quantum Information*, 11:131, 2025.

[14] Bhushan Bonde and Pratik Patil. QProtSim: Quantum-circuit framework for residue-level protein thermodynamic and allosteric simulation. University of Suffolk Open-Source Repository, 2026.

[15] Ramanathan Natesh, Sylva L. U. Schwager, Edward D. Sturrock, and K. Ravi Acharya. Crystal structure of the human angiotensin-converting enzyme–lisinopril complex. *Nature*, 421(6922):551–554, 2003.

[16] Rhyme Quantum Tech. Rhyme beta: High-level symbolic quantum programming framework (`rhyme-beta` v0.1.0b1). Early-Access Software Distribution, 2026.

[17] National Institute of Standards and Technology. Module-lattice-based digital signature standard (FIPS 204: ML-DSA). Technical report, U.S. Department of Commerce, NIST, 2024.

[18] Elham Kashefi. Quantum algorithmic primitives atlas (version 0.4, 27 October 2026). Quantum Software Lab, University of Edinburgh & LIP6, Sorbonne Université / CNRS; originated at Nobel Symposium 210, 2026.

[19] Matthew L. Goh, Martin Larocca, Lukasz Cincio, M. Cerezo, and Frédéric Sauvage. Lie-algebraic classical simulations for quantum computing. *Physical Review Research*, 7:033266, 2025.

[20] Sarah Montgomery-Taylor, Hamsa Subramaniam, Rob Brisk, Liam Foster, Jack Ching, Andrew Marshall, Jim Taylor, John Hernandez, Jen McKay, Dillon Obika, Nnamdi Ezeanochie, Nira Goren, Tiffany Guo, Rachel Kwon, Arlene Chung, Jacqueline Shreibati, Stella Yeoum, Brett Hatfield, Zelda Gulledge, Ankit Nathu, Ferenc Kazinczi, Yun Liu, Jonathan Krause, Heather Cole-Lewis, Susan Thomas, Megan Jones Bell, Bakul Patel, Joelle Barral, Katherine Chou, Yossi Matias, Dana Gelb Safran, Elizabeth A. McGlynn, and Michael Howell. Towards multi-method evaluation of AI in health: Integrating benchmarks, autoraters, and expert evaluation. *arXiv preprint arXiv:2611.05127*, 2026.

[21] NHS England. DCB0129: Clinical risk management: its application in the manufacture of health IT systems. Technical report, NHS Digital / NHS England, 2018.

[22] NHS England. DCB0160: Clinical risk management: its application in the deployment and use of health IT systems. Technical report, NHS Digital / NHS England, 2018.

[23] Andrew Clegg, Chris Bates, John Young, Ronan Ryan, Linda Nichols, Elizabeth Ann Teale, Mohammed A. Mohammed, John Parry, and Tom Marshall. Development and validation of an electronic frailty index using routine primary care electronic health record data. *Age and Ageing*, 45(3):353–360, 2016.

[24] Rupert A. Payne, Silvia C. Mendonca, Marc N. Elliott, Catherine L. Saunders, Duncan A. Edwards, Martin Marshall, and Martin Roland. Development and validation of the Cambridge Multimorbidity Score. *Canadian Medical Association Journal*, 192(5):E107–E114, 2020.

[25] Feifei Bu, Alexandra Burton, Naomi Launders, Amelia E. Taylor, Alvin Richards-Belle, Stephanie Tierney, David Osborn, and Daisy Fancourt. National roll-out of social prescribing in England's primary care system: a longitudinal observational study using Clinical Practice Research Datalink data. *The Lancet Public Health*, 10(11):e903–e911, 2025.

[26] Quantinuum and SoftBank Corp. Quantum computing enterprise roadmap and fault-tolerant architecture white paper. Technical report, Quantinuum Ltd., 2026. Industry Strategic White Paper, July 2026.

[27] Jacob S. Nelson, Andrew J. Landahl, and Andrew D. Baczewski. Below-threshold execution of fault-tolerant ripple-carry addition on a trapped-ion quantum computer. *arXiv preprint arXiv:2610.14202*, 2026.

[28] Timothy Proctor, Oliver Hart, Oliver Widzowski Maupin, Matthew Girling, Daniel Hothem, Daniel Mills, Jordan Hines, Karl Mayer, Jacob S. Nelson, Tyler LeBlond, Zohim Chandani, Diego Forlivesi, Piper C. Wysocki, Boldizsár Poór, Joan M. Dreiling, Annie Park, Adam P. Reed, Brian Estey, Cameron Foltz, Akhil Isanaka, M. S. Allman, Michael Mills, Maxwell D. Urmey, Peter E. Siegfried, Audrey Faricy, Jin-Sung Kim, Cristina Cîrstoiu, Andrew D. Baczewski, Charles H. Baldwin, and Robin Blume-Kohout. Benchmarking the computational power of quantum computers. *arXiv preprint arXiv:2609.12146*, 2026.

[29] Austin G. Fowler, Matteo Mariantoni, John M. Martinis, and Andrew N. Cleland. Surface codes: Towards practical large-scale quantum computation. *Physical Review A*, 86(3):032324, 2012.

[30] Daniel Litinski. A game of surface codes: Large-scale quantum computing with lattice surgery. *Quantum*, 3:128, 2019.

[31] András Gilyén, Yuan Su, Guang Hao Low, and Nathan Wiebe. Quantum singular value transformation and beyond: exponential improvements for quantum matrix arithmetics. *Proceedings of the 51st Annual ACM SIGACT Symposium on Theory of Computing (STOC)*, pages 193–204, 2019.

[32] Gilles Brassard, Peter Hoyer, Michele Mosca, and Alain Tapp. Quantum amplitude amplification and estimation. *Contemporary Mathematics*, 305:53–74, 2002.

[33] Ashley Montanaro. Quantum speedup of branch-and-bound algorithms. *Physical Review Research*, 2(1):013056, 2020.

[34] Jean-Luc Fauchère and Vladimir Pliška. Hydrophobic parameters π of amino-acid side chains from the partitioning of N-acetyl-amino-acid amides. *European Journal of Medicinal Chemistry*, 18(4):369–375, 1983.

[35] Matthew G. Durrant, Nicholas T. Perry, James J. Pai, Aditya R. Jangid, Januka S. Athukoralage, Masahiro Hiraizumi, John P. McSpedon, April Pawluk, Hiroshi Nishimasu, Silvana Konermann, and Patrick D. Hsu. Bridge RNAs direct programmable recombination of target and donor DNA. *Nature*, 630(8018):984–993, 2024.

[36] Marco Cerezo, Akira Sone, Tyler Volkoff, Lukasz Cincio, and Patrick J. Coles. Cost function dependent barren plateaus in shallow parametrized quantum circuits. *Nature Communications*, 12(1):1791, 2021.